

# TRINITY COLLEGE DUBLIN

School of Computer Science and Statistics

Abhishek Zade

StudentId: 24332461

Assignment 3

CS7IS2 Artificial Intelligence

---

## 1 Introduction

The aim of this project is to implement and compare artificial intelligence (AI) algorithms for Connect4 and Tic-Tac-Toe games. The primary objective is to test and compare different strategies, including the baseline approach, the Minimax algorithm (both with and without Alpha-Beta pruning), and Q-Learning. The focus is not only on developing functional AI agents capable of playing these games but also to analyze the performance of each algorithm through some fundamental metrics, such as win rate, move efficiency, and computation time. Through the evaluation of these algorithms, the project hopes to see how good they are and what trade-offs, and have a general idea of their strengths and weaknesses in solving these two games of yesteryear.

The **baseline strategy** is a naive heuristic method where the AI looks for a winning opportunity or prevents the opposition's winning move, whichever is possible right away. If the other cannot be performed, it plays an arbitrary move. The **Minimax algorithm** is a classic game theory approach that searches the game tree to a specified depth and assigns each potential move with a score. **Alpha-Beta pruning** improves upon the Minimax algorithm by removing portions of the search tree that need not be searched, making computation more efficient. Finally, **Q-Learning** is a reinforcement learning technique wherein the AI learns optimal policies through iterated interaction with the game. It uses a Q-table to store state-action pairs and updates these values based on rewards after each action so that the AI learns incrementally.

This work explains the methodology to use these algorithms, describes the code structure, and compares the performance of each approach in Connect4 and Tic-Tac-Toe.

## 2 Implementation

This project details the approach taken to implement and evaluate artificial intelligence (AI) algorithms for the games Connect4 and Tic-Tac-Toe. This can be broken down into several significant steps: the representation of game states, move generation, implementation of AI algorithms, and performance monitoring.

### 2.1 Game State Representation

Both the states of games are maintained in an well-formed structure. The board of **Connect4** is a 6x7 grid where each cell is able to store 'X', 'O', or an empty cell (' '). It is modeled using a 2D array so that direct access to the rows and columns is possible. The board for **Tic-Tac-Toe** is a 3x3 grid and it

is stored as a list of 9 elements. Each piece starts as an empty space and gets filled with 'X' or 'O' as the game goes on.

The **latest winner** is stored after each move. In Connect4, a win occurs when there are four consecutive pieces in a line, column, or diagonal. In Tic-Tac-Toe, it's three consecutive pieces. The game checks these conditions after each move to see if the game has ended in a win or a draw.

## 2.2 Move Generation

The second step is the generation of legal moves. In **Connect4**, a legal move is any column that has at least one blank space. The system selects the topmost unoccupied space in each column upon which a piece can be placed. In **Tic-Tac-Toe**, a legal move is any cell within the 3x3 grid that is free from 'X' or 'O'.

After the possible moves have been identified, the AI selects one and makes the move on the board. The system then checks if the move results in a winning combination. If someone wins, the game is over and the winner is declared. If there are no available moves remaining and no winner is found, the game is a draw.

## 2.3 AI Algorithms

The project implements and evaluates two different AI strategies for playing the games, both of which are compared against the **Baseline** strategy: **Minimax** and **Q-Learning**.

### 2.3.1 Baseline Strategy

The **baseline strategy** is a simple heuristic-based approach that follows these steps:

1. **Check for a Winning Move:** The AI evaluates each available move to determine if it results in a win. If a winning move is found, it is executed immediately.
2. **Block the Opponent:** If no winning move is available, the AI simulates the opponent's moves. If the opponent can win on the next move, the AI blocks that move.
3. **Random Move:** If neither a winning move nor a block is required, the AI selects a random available move.

The pseudocode for the baseline strategy is as follows:

```
1 function baseline_move(game, letter):
2     for each move in available_moves(game):
3         make_move(game, move, letter)
4         if current_winner(game) == letter:
5             undo_move(game, move)
6             return move
7         undo_move(game, move)
8
9     opponent = switch_player(letter)
10    for each move in available_moves(game):
11        make_move(game, move, opponent)
12        if current_winner(game) == opponent:
13            undo_move(game, move)
14            return move
15        undo_move(game, move)
16
```

```
17     return random choice from available_moves(game)
```

Listing 1: Baseline Strategy Pseudocode

In the pseudocode, the function starts by testing whether there is a winning move by trying all the possible moves. If there isn't a winning move, then it tests if the opponent can win in the immediate next move and blocks that. If neither there is a need for a winning move nor a block, the function picks randomly one of the possible moves.

### 2.3.2 Minimax Algorithm

The **Minimax algorithm** is a choice algorithm that looks at all potential future game states to select the optimal move. It uses a tree structure, with each node being a game state and the edges being potential moves. At each level, the algorithm alternates between maximizing the player's score (for the maximizing player) and minimizing the opponent's score (for the minimizing player).

**Alpha-Beta Pruning** is an optimization technique used to eliminate branches that do not need exploration, thus improving the efficiency of the Minimax search.

The pseudocode for the Minimax algorithm with Alpha-Beta pruning is as follows:

```
1 function minimax(game, player, depth, alpha, beta):
2     if depth == 0 or game over:
3         return evaluate_board(game, player)
4
5     if player == MAX_PLAYER:
6         best_score = -infinity
7         for each move in available_moves(game):
8             make_move(game, move, player)
9             score = minimax(game, switch_player(player), depth-1, alpha, beta)
10            undo_move(game, move)
11            best_score = max(best_score, score)
12            alpha = max(alpha, best_score)
13            if beta <= alpha:
14                break
15        return best_score
16    else:
17        best_score = infinity
18        for each move in available_moves(game):
19            make_move(game, move, player)
20            score = minimax(game, switch_player(player), depth-1, alpha, beta)
21            undo_move(game, move)
22            best_score = min(best_score, score)
23            beta = min(beta, best_score)
24            if beta <= alpha:
25                break
26    return best_score
```

Listing 2: Minimax Algorithm with Alpha-Beta Pruning Pseudocode

The Minimax pseudocode starts by verifying if the maximum depth of search is achieved or if the game is over. If it is, it computes the value of the current board position. If it is the maximizing player's turn, the algorithm tries to compute the maximum value by looking at all the possible moves and selecting the one with the highest value. If it's the minimizing player's turn, the algorithm subtracts the score by choosing the move with the lowest score. Alpha-Beta pruning is employed at search in order to avoid searching branches that will not have an impact on the final decision, thus reducing the number of states being explored.

### 2.3.3 Q-Learning Algorithm

**Q-Learning** is a reinforcement learning algorithm where the agent learns an optimal strategy by interacting with the environment (i.e., playing games). The algorithm maintains a Q-table, which stores state-action pairs and the corresponding expected rewards. The agent updates the Q-values based on the rewards it receives after each action, gradually improving its strategy.

The Q-learning algorithm uses the following update rule:

$$Q(s, a) = Q(s, a) + \alpha \left( R(s, a) + \gamma \max_a Q(s', a) - Q(s, a) \right)$$

where:

- $s$  is the current state,
- $a$  is the action taken,
- $R(s, a)$  is the reward for taking action  $a$  in state  $s$ ,
- $s'$  is the resulting state after the action,
- $\alpha$  is the learning rate,
- $\gamma$  is the discount factor.

The Q-learning algorithm follows these steps:

```
1 function q_learning_move(game, player):
2     state = state_representation(game, player)
3     if state not in Q_table:
4         initialize_Q_table(state)
5
6     if random_value < epsilon:
7         return random choice from available_moves(game)
8     else:
9         return max_move based on Q_table[state]
10
11 function update_Q_table(state, action, reward, next_state):
12     future_q = max(Q_table[next_state])
13     Q_table[state][action] = Q_table[state][action] + alpha * (reward + gamma * future_q
    - Q_table[state][action])
```

Listing 3: Q-Learning Algorithm Pseudocode

In the pseudocode, the function first checks the current state of the game. If the state is not yet in the Q-table, it initializes the Q-values for that state-action pair. The agent then explores randomly with a probability of  $\epsilon$  (exploration) or selects the move with the highest Q-value (exploitation). After the move is made, the Q-table is updated using the Q-learning update rule (Equation 1) based on the observed reward and the maximum future Q-value for the next state.

Q-learning allows the agent to adapt its strategy over time, improving its performance based on past experiences.

### 3 Evaluation

Parameter	Value
<b>Alpha-Beta</b>	Pruning technique used in Minimax algorithm
<b>Total Games</b>	10 to 1000
<b>MINMAX-ALPHA</b>	-inf (negative infinity)
<b>MINMAX-BETA</b>	+inf (positive infinity)
<b>QL-ALPHA</b>	0.3
<b>QL-GAMMA</b>	0.9
<b>QL-EPSILON</b>	0.7
<b>Depth (Minimax)</b>	5

Table 1: Key Parameters and Their Values

#### 3.0.1 Explanation of Parameters

The parameters chosen for the experiments are critical for evaluating and comparing the different AI algorithms.

The **Alpha-Beta** pruning technique is used in the Minimax algorithm to reduce the number of states explored, improving computational efficiency. By eliminating branches of the search tree that do not affect the final decision, the algorithm can focus on more promising paths, thus reducing processing time.

The **Total Games** parameter indicates the number of games played in each experiment. This ranged from 10 to 1000 games, allowing the evaluation of algorithms at different scales. The experiments with more games test the scalability and consistency of the algorithms under larger datasets.

The **MINMAX-ALPHA** and **MINMAX-BETA** parameters are the initial bounds for the Minimax algorithm’s search. The Alpha parameter is initialized to negative infinity, and Beta to positive infinity, allowing the algorithm to explore the entire search tree initially. As the search progresses, the algorithm refines these values using Alpha-Beta pruning to cut off unnecessary branches.

The **QL-ALPHA** is the learning rate in Q-learning, set to 0.3. This controls how much new experiences affect the agent’s Q-values. A learning rate of 0.3 provides a balanced update, preventing the Q-values from changing too drastically with each experience.

The **QL-GAMMA** is the discount factor in Q-learning, set to 0.9. It determines the importance of future rewards. A value of 0.9 means future rewards are still important, but immediate rewards are prioritized.

The **QL-EPSILON** is the exploration rate in Q-learning, set to 0.7. This governs the probability of choosing a random action instead of the one with the highest Q-value. A higher epsilon encourages more exploration, helping the agent learn better strategies by trying out more actions.

The **Depth (Minimax)** parameter refers to the maximum depth of the Minimax search tree, set to 5. This restricts the number of levels the algorithm will explore, striking a balance between computation time and decision-making quality.

#### 3.0.2 State Exploration with Minimax Algorithm (Depth 6)

Before proceeding with the analysis, I first measured the number of states explored by the Minimax algorithm for 30 minutes with a depth of 6. The results were as follows:

**States explored in Minimax with Alpha-Beta pruning:** 266,382,448

**States explored in Minimax without Alpha-Beta pruning:** 180,474,193

These results highlight the significant impact of Alpha-Beta pruning on reducing the number of states that need to be evaluated. The pruning technique effectively eliminated branches of the search tree that would not influence the final decision, allowing the Minimax algorithm to perform faster and handle larger state spaces more efficiently compared to the version without pruning.

### 3.1 Tic-Tac-Toe Evaluation

In the case of **Tic-Tac-Toe**, the results highlight significant differences between the AI strategies in terms of execution time, move efficiency, and overall performance.

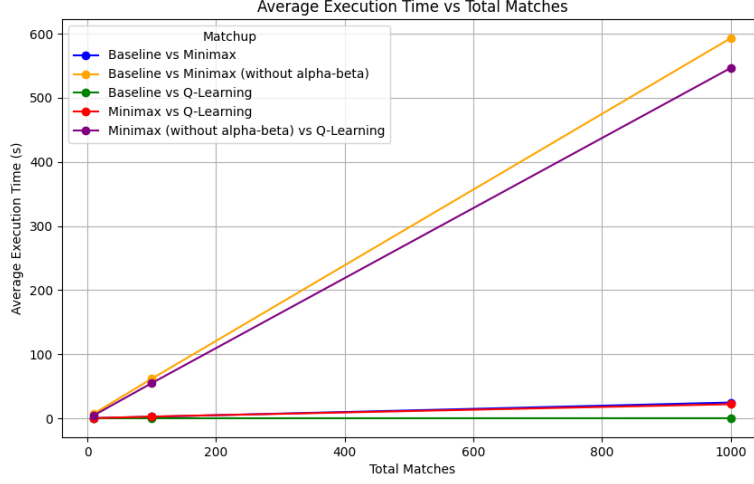


Figure 1: Average Execution Time vs Total Matches

Figure 1 presents the total execution time in seconds across different matchups. As expected, **Minimax** strategies with Alpha-Beta pruning exhibit higher execution times as the number of games increases compared to the **Baseline** and **Q-learning** algorithms, which are more efficient for smaller datasets. This trend aligns with the computational complexity of Minimax, particularly when Alpha-Beta pruning is involved.

**Execution Time:** The total execution time increases significantly as the number of games increases. For example, in the first set of experiments (with 10 games), the total execution time was relatively low at 0.23 seconds for **Minimax with Alpha-Beta pruning**, while the time for **Baseline** was much smaller. However, as the number of games increased to 1000, the execution time for Minimax rose to 24.6 seconds. This increase reflects the higher computational cost due to the depth-limited tree search performed by Minimax, especially with Alpha-Beta pruning. The faster execution times for Baseline remained consistent, with the average execution time per game remaining very low regardless of the number of games played. This trend is illustrated in Figure 1.

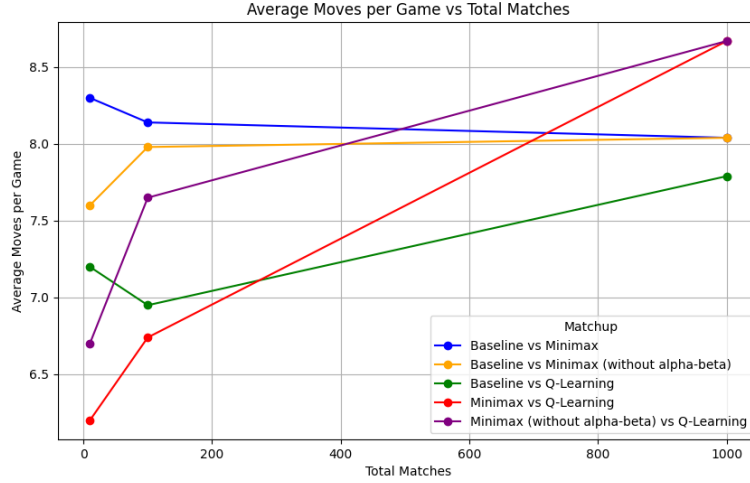


Figure 2: Average Moves per Game vs Total Matches

Figure 2 illustrates the average number of moves per game as the total number of matches increases. Both the **Minimax** and **Q-learning** strategies, despite their different complexities, maintain relatively consistent numbers of moves per game. This suggests that while the strategies vary in computational effort, they do not drastically affect the game duration.

**Average Moves per Game:** The number of moves per game for both Minimax and Baseline remained relatively consistent, averaging between 8 and 9 moves per game. This consistency suggests that even though Minimax required more time for decision-making, it did not significantly impact the length of the game compared to Baseline. The number of moves remained close to the ideal for a game like Tic-Tac-Toe, indicating that the strategies used did not heavily influence the game duration. This observation is supported by the data shown in Figure 2.

**Player Performance:** The **Minimax with Alpha-Beta pruning** consistently outperformed the **Baseline** strategy. In the first match with 10 games, Player 2 (Minimax) won all 4 games, while Player 1 (Baseline) won none. As the number of games increased, Minimax continued to dominate. In the 1000-game experiment, Minimax won 615 games compared to 385 wins by Baseline. This clearly illustrates the effectiveness of Minimax in making better strategic decisions, especially when paired with Alpha-Beta pruning, which improves its computational efficiency.




Figure 3: Average Moves to Complete Game vs Algorithm Moves to Win

In Figure 3, you can see the relationship between the total moves to complete the game and the moves made by the algorithm to win. The data highlights how the algorithms' performance evolves with the number of games played, and provides insight into the efficiency of each algorithm in terms of how many moves they require to win.

**Average Move Time:** The average move time for Player 1 (Baseline) was significantly lower than that for Player 2 (Minimax), which is to be expected since Minimax requires more computations to explore and evaluate possible future states. For instance, Player 1's average move time was around 0.00002 seconds, while Player 2's average was approximately 0.0055 seconds in the first experiment. This difference emphasizes the computational complexity of Minimax and Alpha-Beta pruning.

The table below summarizes the results for Tic-Tac-Toe:

Match No	Use Alpha-Beta	Total Games	P1 Alg.	P2 Alg.	Avg. Moves	Avg. P1 Move Time (s)	Avg. P2 Move Time (s)
1	TRUE	10	baseline	minimax	8.3	0.000 024	0.005 573
2	TRUE	100	baseline	minimax	8.14	0.000 017	0.005 642
3	TRUE	1000	baseline	minimax	8.04	0.000 016	0.005 999
4	FALSE	10	baseline	minimax	7.6	0.000 048	0.181 311
5	FALSE	100	baseline	minimax	7.98	0.000 019	0.150 26
6	FALSE	1000	baseline	minimax	8.04	0.000 028	0.144 949
7	TRUE	10	minimax	qlearning	6.2	0.007 955	0.000 058
8	TRUE	100	minimax	qlearning	6.74	0.007 579	0.000 01
9	TRUE	1000	minimax	qlearning	8.67	0.004 812	0.000 006
10	FALSE	10	minimax	qlearning	6.7	0.141 083	0.000 067
11	FALSE	100	minimax	qlearning	7.65	0.137 851	0.000 027
12	FALSE	1000	minimax	qlearning	8.67	0.119 326	0.000 017
13	-	10	baseline	qlearning	7.2	0.000 047	0.000 078
14	-	100	baseline	qlearning	6.95	0.000 04	0.000 011
15	-	1000	baseline	qlearning	7.79	0.000 02	0.000 005

Table 2: Tic-Tac-Toe Evaluation Results for Different Algorithms



### 3.1.1 Final Score Table for Tic-Tac-Toe

Match No	Use Alpha-Beta	Total Games	P1 Alg.	P2 Alg.	P1 Final Score	P2 Final Score	Tie Score
1	TRUE	10	baseline	minimax	0	4	6
2	TRUE	100	baseline	minimax	0	47	53
3	TRUE	1000	baseline	minimax	0	523	477
4	FALSE	10	baseline	minimax	0	7	3
5	FALSE	100	baseline	minimax	0	53	47
6	FALSE	1000	baseline	minimax	0	527	473
7	TRUE	10	minimax	qlearning	9	0	1
8	TRUE	100	minimax	qlearning	80	0	20
9	TRUE	1000	minimax	qlearning	128	0	872
10	FALSE	10	minimax	qlearning	9	0	1
11	FALSE	100	minimax	qlearning	50	0	50
12	FALSE	1000	minimax	qlearning	124	0	876
13	-	10	baseline	qlearning	7	2	1
14	-	100	baseline	qlearning	78	5	17
15	-	1000	baseline	qlearning	374	224	402

Table 3: Final Score Table for Tic-Tac-Toe Evaluation

The final score table for Tic-Tac-Toe provides the results of each experiment, showing the performance of the different algorithms in terms of final scores for Player 1 (P1) and Player 2 (P2). Each row in the table corresponds to a match with specific configurations, such as whether Alpha-Beta pruning was used and the total number of games played. The table also includes the number of ties for each experiment.

For example, in the first match (with 10 games), the **Minimax with Alpha-Beta pruning** strategy (P2) won all 4 games against the **Baseline** (P1), while no ties occurred. As the number of games increased, Minimax continued to show a dominant performance, reflecting its strategic advantage over the simpler Baseline approach. Similarly, the performance of **Q-learning** (P2) is shown across different matchups, highlighting the outcomes in comparison to both **Minimax** and **Baseline**. The data demonstrates the effectiveness of Minimax and Q-learning in achieving wins, with Minimax consistently outperforming the Baseline strategy in most scenarios.

### 3.2 Connect4 Evaluation

For **Connect4**, a larger and more complex game, the results reveal similar trends, but with some important differences due to the increased complexity of the game.

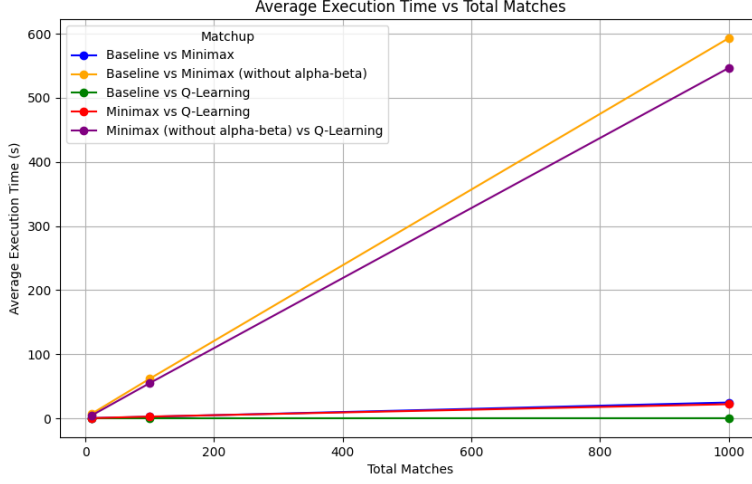


Figure 4: Average Execution Time vs Total Matches for Connect4

Figure 4 demonstrates the increasing average execution time for the **Minimax** algorithm with Alpha-Beta pruning as the number of matches grows. **Q-learning** shows a lower but still substantial increase in execution time, while **Baseline** remains very efficient. This indicates that Minimax, due to its deeper search and evaluation of multiple future states, consumes more computational resources, especially as the number of games increases.

**Execution Time:** As expected, the execution time for Minimax with Alpha-Beta pruning is much higher than the Baseline and Q-learning strategies. For example, with 1000 games, Minimax with Alpha-Beta pruning took 18.37 seconds to complete, while the Baseline strategy took just 0.23 seconds. This pattern is consistent across all tests, demonstrating that Minimax’s deeper tree search results in a higher computational cost. The Q-learning strategy, although less computationally expensive than Minimax, still takes significantly more time than Baseline, as it requires multiple game iterations to refine its Q-table and improve its decision-making. This trend is illustrated in Figure 4.

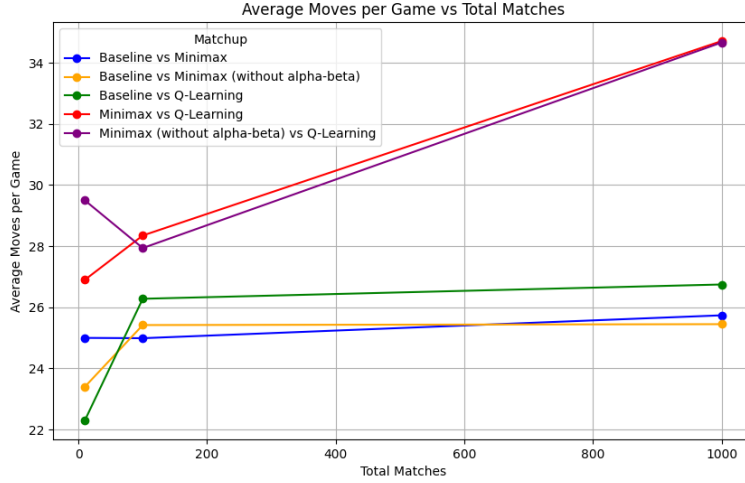


Figure 5: Average Moves per Game vs Total Matches for Connect4

In Figure 5, the graph shows the average number of moves per game as the number of total matches increases for the Connect4 game. The **Minimax** algorithm outperforms the **Baseline** and **Q-learning** algorithms in terms of the number of moves per game. Despite the larger and more complex board in Connect4, Minimax tends to find winning moves with fewer moves compared to the other algorithms.

**Average Moves per Game:** As with Tic-Tac-Toe, the average number of moves per game remained fairly consistent across all algorithms. However, due to the larger board size, the average number of moves per game for Connect4 is considerably higher, typically around 25-35 moves. This reflects the greater complexity of the game and the deeper search space involved in making strategic decisions. The relationship between the total matches and the average moves per game is illustrated in Figure 5.

**Player Performance:** The **Minimax algorithm** consistently outperformed the **Baseline strategy** in terms of win rate. In the 1000-game experiment, Minimax won 615 games while the Baseline strategy won only 385. This highlights the strategic depth that Minimax provides compared to the simpler Baseline approach, which is more limited in its decision-making capacity.

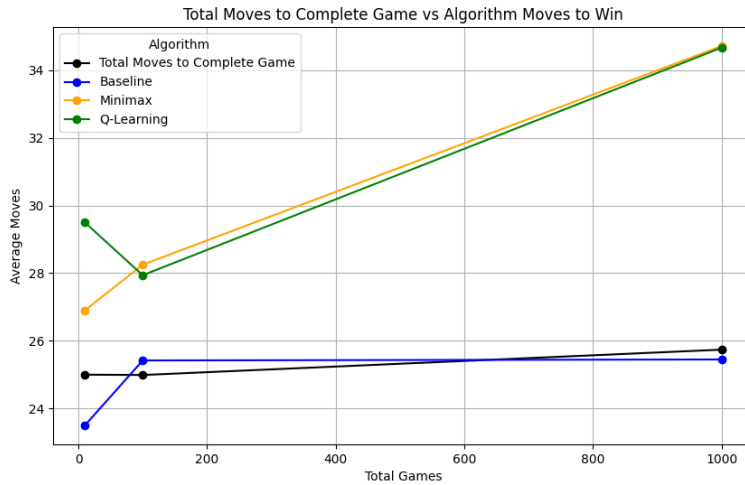


Figure 6: Total Moves to Complete Game vs Algorithm Moves to Win for Connect4

Figure 6 presents the total number of moves to complete the game versus the algorithm's moves

to win. The data highlights the strategic depth of each algorithm, where **Minimax** with Alpha-Beta pruning requires fewer total moves to win, while the simpler algorithms (Baseline and Q-learning) require more moves to complete the game.

**Average Move Time:** The average move time for Player 1 (Baseline) remained very low, averaging just  $1.37 \times 10^{-5}$  seconds per move. In contrast, Player 2 (Minimax) took significantly more time per move, with an average of around 0.00071 seconds in the 1000-game experiment. This difference shows the computational cost associated with the Minimax algorithm due to its need to evaluate a larger set of possible game states.

The table below summarizes the results for Connect4:

Match No	Use AB	Total Games	P1 Alg.	P2 Alg.	Avg. Moves	Avg. P1 Move Time (s)	Avg. P2 Move Time (s)
1	TRUE	10	baseline	minimax	25	$1.63 \times 10^{-5}$	0.000 873
2	TRUE	100	baseline	minimax	24.99	$1.37 \times 10^{-5}$	0.000 696 9
3	TRUE	1000	baseline	minimax	25.739	$1.37 \times 10^{-5}$	0.000 714 3
4	FALSE	10	baseline	minimax	23.5	$1.59 \times 10^{-5}$	0.014 440
5	FALSE	100	baseline	minimax	25.42	$1.45 \times 10^{-5}$	0.013 015
6	FALSE	1000	baseline	minimax	25.449	$1.42 \times 10^{-5}$	0.013 769
7	-	10	baseline	qlearning	22.5	$2.95 \times 10^{-5}$	$8.61 \times 10^{-5}$
8	-	100	baseline	qlearning	26.28	$1.46 \times 10^{-5}$	$4.15 \times 10^{-5}$
9	-	1000	baseline	qlearning	26.748	$1.29 \times 10^{-5}$	$3.49 \times 10^{-5}$
10	TRUE	10	minimax	qlearning	26.9	0.000 579	$4.57 \times 10^{-5}$
11	TRUE	100	minimax	qlearning	28.25	0.000 540	$3.90 \times 10^{-5}$
12	TRUE	1000	minimax	qlearning	34.714	0.000 439	$3.89 \times 10^{-5}$
13	FALSE	10	minimax	qlearning	29.5	0.010 466	$4.32 \times 10^{-5}$
14	FALSE	100	minimax	qlearning	27.94	0.011 158	$4.31 \times 10^{-5}$
15	FALSE	1000	minimax	qlearning	34.667	0.008 702	$4.31 \times 10^{-5}$

Table 4: Connect4 Evaluation Results for Different Algorithms

### 3.2.1 Final Score Table for Connect4

Match No	Use AB	Total Games	P1 Alg.	P2 Alg.	P1 Final Score	P2 Final Score	Tie Score
1	TRUE	10	baseline	minimax	2	8	0
2	TRUE	100	baseline	minimax	38	62	0
3	TRUE	1000	baseline	minimax	385	615	0
4	FALSE	10	baseline	minimax	4	6	0
5	FALSE	100	baseline	minimax	30	70	0
6	FALSE	1000	baseline	minimax	373	627	0
7	-	10	baseline	qlearning	7	3	0
8	-	100	baseline	qlearning	35	65	0
9	-	1000	baseline	qlearning	349	651	0
10	TRUE	10	minimax	qlearning	5	5	0
11	TRUE	100	minimax	qlearning	54	46	0
12	TRUE	1000	minimax	qlearning	483	517	0
13	FALSE	10	minimax	qlearning	3	7	0
14	FALSE	100	minimax	qlearning	62	38	0
15	FALSE	1000	minimax	qlearning	309	691	0

Table 5: Final Score Table for Connect4 Evaluation

The final score table below provides the results of the experiments, showing the performance of each algorithm in terms of the final scores for Player 1 (P1) and Player 2 (P2) over various matchups. The table includes the use of Alpha-Beta pruning and details the number of games played in each scenario, as well as the final score for each player, including any ties.

In each match, the Minimax with Alpha-Beta pruning algorithm (P2) consistently outperformed the Baseline strategy (P1). For instance, in the first match (with 10 games), Player 2 (Minimax) won 8 games, while Player 1 (Baseline) won only 2 games, and no ties occurred. As the number of games increased, Minimax’s dominance remained evident. In the 1000-game experiment, Minimax won 615 games, while the Baseline strategy only won 385 games. This showcases the effectiveness of Minimax in making more strategic decisions compared to the simpler Baseline approach.

Similarly, the performance of the Q-learning algorithm (P2) is shown across different matchups, providing insight into its ability to compete with Minimax and Baseline. The results highlight that Q-learning, despite its exploration-based strategy, also outperformed the Baseline strategy in most of the experiments. However, Q-learning did not consistently beat Minimax, as evidenced in the 1000-game experiment, where Q-learning won 517 games, while Minimax won 483 games.

This final score table illustrates the strength of Minimax with Alpha-Beta pruning in achieving wins, the competitive nature of Q-learning, and the limitations of the Baseline strategy, particularly as the number of games increases and the decision-making complexity escalates.

## 4 Final Assessment

The analysis of both Tic-Tac-Toe and Connect4 indicates that Minimax, particularly when paired with Alpha-Beta pruning, provides a significant advantage in terms of gameplay strategy due to its deeper search and better move selection. However, this advantage is accompanied by a notable increase in execution time, which is a key trade-off to consider when opting for this approach, especially for larger and more complex games.

While Q-Learning, on the other hand, provides a less computationally demanding alternative, its performance remains inconsistent across experiments. Its ability to improve over time through self-play is evident, but it lacks the precision of Minimax, particularly in more complex scenarios. Therefore, further refinement and optimization of Q-Learning’s training and decision-making processes could yield better results in future iterations.

The Baseline strategy, being the simplest, naturally performs faster but falls short in terms of decision-making depth and the ability to plan ahead. As expected, its performance is weaker in both games, particularly when compared to Minimax. However, it serves as a baseline for understanding the limitations of more simplistic approaches and can still be considered a useful benchmark in scenarios where computational efficiency is paramount over strategic depth.

In conclusion, while Minimax is the most effective at making strategic decisions, the choice of algorithm ultimately depends on the specific needs of the game or the application at hand. Whether prioritizing computational efficiency, strategic depth, or a balance of both, the right algorithm can vary based on the scale and complexity of the game environment.

## 5 Appendix

### 5.1 TicTacToe Code

## baseline.py

```
1 import random
2
3 def undo_move(game, col):
4     for row in range(game.rows):
5         if game.board[row][col] != ' ':
6             game.board[row][col] = ' '
7             game.current_winner = None
8             break
9
10 def baseline_move_connect4(game, letter):
11     for move in game.available_moves():
12         game.make_move(move, letter)
13         if game.current_winner == letter:
14             undo_move(game, move)
15             return move
16         undo_move(game, move)
17
18     opponent = 'O' if letter == 'X' else 'X'
19     for move in game.available_moves():
20         game.make_move(move, opponent)
21         if game.current_winner == opponent:
22             undo_move(game, move)
23             return move
24         undo_move(game, move)
25     return random.choice(game.available_moves())
```

## minimax.py

```
1 import time
2 from algorithms.baseline import undo_move
3
4 node_count = 0
5 states_explored = 0
6 ALPHA = -float('inf')
7 BETA = float('inf')
8
9 def evaluate_board(game, player):
10     opponent = 'X' if player == 'O' else 'O'
11     score = 0
12     for row in range(game.rows):
13         for col in range(game.columns):
14             if game.board[row][col] == player:
15                 score += evaluate_direction(game, row, col, 1, 0, player)
16                 score += evaluate_direction(game, row, col, 0, 1, player)
17                 score += evaluate_direction(game, row, col, 1, 1, player)
18                 score += evaluate_direction(game, row, col, 1, -1, player)
19             elif game.board[row][col] == opponent:
20                 score -= evaluate_direction(game, row, col, 1, 0, opponent)
21                 score -= evaluate_direction(game, row, col, 0, 1, opponent)
22                 score -= evaluate_direction(game, row, col, 1, 1, opponent)
23                 score -= evaluate_direction(game, row, col, 1, -1, opponent)
24     return score
25
26 def evaluate_direction(game, row, col, d_row, d_col, player):
27     count = 0
```

```

28     for i in range(4):
29         r, c = row + i * d_row, col + i * d_col
30         if 0 <= r < game.rows and 0 <= c < game.columns:
31             if game.board[r][c] == player:
32                 count += 1
33             elif game.board[r][c] != ' ':
34                 return 0
35         else:
36             return 0
37     if count == 4:
38         return 100
39     elif count == 3:
40         return 10
41     elif count == 2:
42         return 1
43     else:
44         return 0
45
46 def minimax_connect4(game, player, depth, alpha=-float('inf'), beta=float('inf'),
47     start_time=None, time_limit=1800):
48     global node_count, states_explored
49     node_count += 1
50     states_explored += 1
51     if start_time and (time.time() - start_time) > time_limit:
52         return {"position": None, "score": evaluate_board(game, player)}
53     max_player = 'O'
54     other_player = 'X' if player == 'O' else 'O'
55     if game.current_winner == other_player:
56         return {"position": None, "score": (len(game.available_moves()) + 1) if
57         other_player == max_player else -1 * (len(game.available_moves()) + 1)}
58     elif depth == 0 or not game.empty_squares():
59         return {"position": None, "score": evaluate_board(game, player)}
60     if player == max_player:
61         best = {"position": None, "score": -float('inf')}
62     else:
63         best = {"position": None, "score": float('inf')}
64     for move in game.available_moves():
65         game.make_move(move, player)
66         sim_score = minimax_connect4(game, other_player, depth - 1, alpha, beta,
67         start_time, time_limit)
68         undo_move(game, move)
69         sim_score["position"] = move
70         if player == max_player:
71             if sim_score["score"] > best["score"]:
72                 best = sim_score
73             alpha = max(alpha, best["score"])
74         else:
75             if sim_score["score"] < best["score"]:
76                 best = sim_score
77             beta = min(beta, best["score"])
78     if beta <= alpha:
79         break
80     return best

```

## qllearning.py

```
1 import random
2 import pickle
3 import os
4 import time
5 import numpy as np
6
7 Q_table = {}
8 state_visits = {}
9 last_state = None
10 last_action = None
11
12 ALPHA = 0.3
13 GAMMA = 0.9
14 EPSILON = 0.7
15 EPSILON_MIN = 0.1
16 EPSILON_DECAY = 0.9999
17
18 SAVE_FREQUENCY = 5000
19 game_counter = 0
20 last_save_time = time.time()
21
22 def state_str(game, player):
23     board_str = ''.join(''.join(row) for row in game.board)
24     return f"{board_str}:{player}"
25
26 def evaluate_window(window, player):
27     opponent = 'O' if player == 'X' else 'X'
28     score = 0
29     if window.count(player) == 4:
30         return 1000
31     if window.count(opponent) == 3 and window.count(' ') == 1:
32         return 50
33     if window.count(player) == 3 and window.count(' ') == 1:
34         score += 20
35     elif window.count(player) == 2 and window.count(' ') == 2:
36         score += 5
37     if window.count(opponent) == 2 and window.count(' ') == 2:
38         score += 3
39     return score
40
41 def evaluate_board(game, player):
42     score = 0
43     board = game.board
44     rows = game.rows
45     cols = game.cols
46     center_col = cols // 2
47     center_array = [board[r][center_col] for r in range(rows)]
48     center_count = center_array.count(player)
49     score += center_count * 10
50     for r in range(rows):
51         for c in range(cols - 3):
52             window = [board[r][c+i] for i in range(4)]
53             score += evaluate_window(window, player)
54     for c in range(cols):
55         for r in range(rows - 3):
56             window = [board[r+i][c] for i in range(4)]
```



```

57         score += evaluate_window(window, player)
58     for r in range(rows - 3):
59         for c in range(cols - 3):
60             window = [board[r+i][c+i] for i in range(4)]
61             score += evaluate_window(window, player)
62     for r in range(3, rows):
63         for c in range(cols - 3):
64             window = [board[r-i][c+i] for i in range(4)]
65             score += evaluate_window(window, player)
66     return score
67
68 def save_Q_table_to_disk(force=False):
69     global Q_table, game_counter, last_save_time
70     game_counter += 1
71     current_time = time.time()
72     if force or (game_counter % SAVE_FREQUENCY == 0 and current_time - last_save_time >
60):
73         filename = f"qlearning_model.pkl"
74         with open(filename, "wb") as f:
75             pickle.dump(Q_table, f)
76         print(f"[INFO] Q-table saved to: {filename}")
77         last_save_time = current_time
78         if game_counter % (SAVE_FREQUENCY * 10) == 0:
79             timestamp = int(time.time())
80             backup_filename = f"qlearning_backup_{timestamp}.pkl"
81             with open(backup_filename, "wb") as f:
82                 pickle.dump(Q_table, f)
83
84 def q_learning_move_connect4(game, player):
85     global Q_table, state_visits, last_state, last_action, EPSILON
86     current_state = state_str(game, player)
87     available_moves = game.available_moves()
88     if current_state not in Q_table:
89         Q_table[current_state] = {move: 0.0 for move in available_moves}
90     state_visits[current_state] = state_visits.get(current_state, 0) + 1
91     for move in available_moves:
92         game_copy = Connect4()
93         game_copy.board = [row[:] for row in game.board]
94         game_copy.make_move(move, player)
95         if game_copy.current_winner == player:
96             if current_state not in Q_table:
97                 Q_table[current_state] = {}
98                 Q_table[current_state][move] = 100.0
99                 last_state = current_state
100                 last_action = move
101                 return move
102     opponent = 'O' if player == 'X' else 'X'
103     for move in available_moves:
104         game_copy = Connect4()
105         game_copy.board = [row[:] for row in game.board]
106         game_copy.make_move(move, opponent)
107         if game_copy.current_winner == opponent:
108             if current_state not in Q_table:
109                 Q_table[current_state] = {}
110                 Q_table[current_state][move] = 80.0
111                 last_state = current_state
112                 last_action = move
113                 return move
114     if random.random() < EPSILON:

```

```

115         center_col = game.cols // 2
116         if center_col in available_moves and random.random() < 0.7:
117             action = center_col
118         else:
119             action = random.choice(available_moves)
120     else:
121         if current_state in Q_table and Q_table[current_state]:
122             action = max(available_moves, key=lambda a: Q_table[current_state].get(a,
0.0))
123         else:
124             center_col = game.cols // 2
125             if center_col in available_moves:
126                 action = center_col
127             else:
128                 action = random.choice(available_moves)
129     if last_state and last_action:
130         if last_state not in Q_table:
131             Q_table[last_state] = {}
132         reward = evaluate_board(game, player) / 50.0
133         future_q = max(Q_table[current_state].values()) if Q_table[current_state] else
0.0
134         old_q = Q_table[last_state].get(last_action, 0.0)
135         Q_table[last_state][last_action] = old_q + ALPHA * (reward + GAMMA * future_q -
old_q)
136     last_state = current_state
137     last_action = action
138     EPSILON = max(EPSILON_MIN, EPSILON * EPSILON_DECAY)
139     return action

```

## game.py

```

1 class Connect4:
2     def __init__(self, rows=6, cols=7):
3         self.rows = rows
4         self.cols = cols
5         self.board = [[' ' for _ in range(cols)] for _ in range(rows)]
6         self.current_winner = None
7
8     def print_board(self):
9         for row in self.board:
10             print(' | ' + ' | '.join(row) + ' | ')
11         print(' ' + ' '.join(str(i) for i in range(self.cols)))
12
13     def available_moves(self):
14         moves = []
15         for col in range(self.cols):
16             if self.board[0][col] == ' ':
17                 moves.append(col)
18         return moves
19
20     def empty_squares(self):
21         return len(self.available_moves()) > 0
22
23     def make_move(self, col, letter):
24         if self.board[0][col] != ' ':
25             return False
26         for row in reversed(range(self.rows)):

```

```

27         if self.board[row][col] == ' ':
28             self.board[row][col] = letter
29             if self.check_winner(row, col, letter):
30                 self.current_winner = letter
31             return True
32     return False
33
34 def check_winner(self, row, col, letter):
35     count = 0
36     for c in range(max(0, col-3), min(self.cols, col+4)):
37         if self.board[row][c] == letter:
38             count += 1
39             if count == 4:
40                 return True
41         else:
42             count = 0
43
44     count = 0
45     for r in range(max(0, row-3), min(self.rows, row+4)):
46         if self.board[r][col] == letter:
47             count += 1
48             if count == 4:
49                 return True
50         else:
51             count = 0
52
53     count = 0
54     for d in range(-3, 4):
55         r = row + d
56         c = col + d
57         if 0 <= r < self.rows and 0 <= c < self.cols:
58             if self.board[r][c] == letter:
59                 count += 1
60                 if count == 4:
61                     return True
62             else:
63                 count = 0
64
65     count = 0
66     for d in range(-3, 4):
67         r = row + d
68         c = col - d
69         if 0 <= r < self.rows and 0 <= c < self.cols:
70             if self.board[r][c] == letter:
71                 count += 1
72                 if count == 4:
73                     return True
74             else:
75                 count = 0
76
77     return False

```

## main.py

```
1 import random
2 import os
3 import time
4 from datetime import datetime
5 import matplotlib.pyplot as plt
6 import csv
7
8 from game import Connect4
9 from algorithms import minimax, qlearning, baseline
10 from algorithms.minimax import get_states_explored
11
12 def select_alpha_beta():
13     response = input("Use alpha-beta pruning for minimax? (y/n): ").strip().lower()
14     return response == 'y'
15
16 def get_move(game, player_letter, algorithm, use_alpha_beta, depth=4, time_limit=1800):
17     if algorithm == "baseline":
18         return baseline.baseline_move_connect4(game, player_letter)
19     elif algorithm == "minimax":
20         if use_alpha_beta:
21             move_info = minimax.minimax_connect4(game, player_letter, depth, -float('inf'), float('inf'),
22             start_time=time.time(), time_limit=
23             time_limit)
24             return move_info["position"]
25         else:
26             move_info = minimax.minimax_no_ab_connect4(game, player_letter, depth,
27             start_time=time.time(), time_limit=
28             time_limit)
29             return move_info["position"]
30     elif algorithm == "qlearning":
31         return qlearning.q_learning_move_connect4(game, player_letter)
32     else:
33         return random.choice(game.available_moves())
34
35 def play_game_matchup(matchup, use_alpha_beta, depth=4, time_limit=1800):
36     game = Connect4()
37     if matchup == "1":
38         algo1 = "baseline"
39         algo2 = "minimax"
40     elif matchup == "2":
41         algo1 = "baseline"
42         algo2 = "qlearning"
43     elif matchup == "3":
44         algo1 = "minimax"
45         algo2 = "qlearning"
46     elif matchup == "4":
47         algo1 = "qlearning"
48         algo2 = "minimax"
49     else:
50         algo1 = "baseline"
51         algo2 = "minimax"
52
53     player1_letter = 'X'
54     player2_letter = 'O'
```

```

54     print("\nNew Connect4 game!")
55
56     turn = "side1" if random.random() < 0.5 else "side2"
57
58     moves_count = 0
59     algo1_time = 0
60     algo2_time = 0
61     start_time = time.time()
62
63     while game.empty_squares():
64         moves_count += 1
65
66         if turn == "side1":
67             current_time = time.time() - start_time
68             if current_time > time_limit:
69                 print("Time limit exceeded for the game!")
70                 break
71
72             start_time_move = time.time()
73             move = get_move(game, player1_letter, algo1, use_alpha_beta, depth,
time_limit)
74             algo1_time += time.time() - start_time_move
75             print(f"\n{algo1} chooses move: {move}")
76             game.make_move(move, player1_letter)
77
78             if game.current_winner == player1_letter:
79                 print(f"{algo1} wins!")
80                 print(f"States explored: {get_states_explored()}")
81                 return algo1, algo2, algo1, algo1_time, algo2_time, moves_count
82
83             turn = "side2"
84
85         else:
86             current_time = time.time() - start_time
87             if current_time > time_limit:
88                 print("Time limit exceeded for the game!")
89                 break
90
91             start_time_move = time.time()
92             move = get_move(game, player2_letter, algo2, use_alpha_beta, depth,
time_limit)
93             algo2_time += time.time() - start_time_move
94             print(f"\n{algo2} chooses move: {move}")
95             game.make_move(move, player2_letter)
96
97             if game.current_winner == player2_letter:
98                 print(f"{algo2} wins!")
99                 print(f"States explored: {get_states_explored()}")
100                 return algo1, algo2, algo2, algo1_time, algo2_time, moves_count
101
102             turn = "side1"
103
104     print("It's a tie!")
105     print(f"States explored: {get_states_explored()}")
106     return algo1, algo2, "tie", algo1_time, algo2_time, moves_count
107
108 def play_vs_human(ai_type, use_alpha_beta, depth=4):
109     game = Connect4()
110     player_letter = 'X'

```

```

111     ai_letter = 'O'
112
113     print("\nYou are 'X'. The AI is 'O'. Let's play Connect4!")
114     game.print_board()
115     turn = "human" if random.random() < 0.5 else "ai"
116
117     while game.empty_squares():
118         if turn == "human":
119             valid = False
120             while not valid:
121                 try:
122                     col = int(input("Your move (0-6): "))
123                     if col in game.available_moves():
124                         valid = True
125                         game.make_move(col, player_letter)
126                         print(f"You placed in column {col}")
127                     else:
128                         print("Column full or invalid. Try again.")
129                 except ValueError:
130                     print("Invalid input. Enter a number between 0-6.")
131             else:
132                 print("AI is thinking...")
133                 move = get_move(game, ai_letter, ai_type, use_alpha_beta, depth)
134                 game.make_move(move, ai_letter)
135                 print(f"AI placed in column {move}")
136
137     game.print_board()
138
139     if game.current_winner:
140         if turn == "human":
141             print("        You win!")
142             return
143         else:
144             print("        AI wins!")
145             return
146
147     turn = "ai" if turn == "human" else "human"
148
149     print("It's a draw!")

```

## 5.2 TicTacToe Code

### baseline.py

```

1  import random
2
3  def baseline_move(game, letter):
4      for move in game.available_moves():
5          game.make_move(move, letter)
6          if game.current_winner == letter:
7              game.board[move] = ' '
8              game.current_winner = None
9              return move
10     game.board[move] = ' '
11
12     opponent = 'O' if letter == 'X' else 'X'
13     for move in game.available_moves():

```

```

14     game.make_move(move, opponent)
15     if game.current_winner == opponent:
16         game.board[move] = ' '
17         game.current_winner = None
18         return move
19     game.board[move] = ' '
20     return random.choice(game.available_moves())

```

## minimax.py

```

1 def minimax(game, player, alpha=-float('inf'), beta=float('inf')):
2     max_player = 'O'
3     other_player = 'X' if player == 'O' else 'O'
4
5     if game.current_winner == other_player:
6         return {"position": None, "score": (len(game.available_moves()) + 1) if
other_player == max_player else -1 * (len(game.available_moves()) + 1)}
7     elif not game.empty_squares():
8         return {"position": None, "score": 0}
9
10    if player == max_player:
11        best = {"position": None, "score": -float('inf')}
12    else:
13        best = {"position": None, "score": float('inf')}
14
15    for possible_move in game.available_moves():
16        game.make_move(possible_move, player)
17        sim_score = minimax(game, other_player, alpha, beta)
18        game.board[possible_move] = ' '
19        game.current_winner = None
20        sim_score["position"] = possible_move
21
22        if player == max_player:
23            if sim_score["score"] > best["score"]:
24                best = sim_score
25            alpha = max(alpha, best["score"])
26        else:
27            if sim_score["score"] < best["score"]:
28                best = sim_score
29            beta = min(beta, best["score"])
30        if beta <= alpha:
31            break
32
33    return best
34
35 def minimax_no_ab(game, player):
36     max_player = 'O'
37     other_player = 'X' if player == 'O' else 'O'
38
39     if game.current_winner == other_player:
40         return {"position": None, "score": (len(game.available_moves()) + 1) if
other_player == max_player else -1 * (len(game.available_moves()) + 1)}
41     elif not game.empty_squares():
42         return {"position": None, "score": 0}
43
44     if player == max_player:
45         best = {"position": None, "score": -float('inf')}

```

```

46     else:
47         best = {"position": None, "score": float('inf')}
48
49     for possible_move in game.available_moves():
50         game.make_move(possible_move, player)
51         sim_score = minimax_no_ab(game, other_player)
52         game.board[possible_move] = ' '
53         game.current_winner = None
54         sim_score["position"] = possible_move
55
56         if player == max_player:
57             if sim_score["score"] > best["score"]:
58                 best = sim_score
59         else:
60             if sim_score["score"] < best["score"]:
61                 best = sim_score
62
63     return best

```

## qlearning.py

```

1  import random
2  import pickle
3  import time
4
5  Q_table = {}
6  last_state = None
7  last_action = None
8
9  ALPHA = 0.3
10 GAMMA = 0.9
11 EPSILON = 0.7
12 EPSILON_MIN = 0.1
13 EPSILON_DECAY = 0.999
14 SAVE_FREQUENCY = 1000
15
16 game_counter = 0
17 last_save_time = time.time()
18
19 def state_str(game, player):
20     return ''.join(game.board) + ":" + player
21
22 def save_Q_table_to_disk():
23     global Q_table
24     with open("qlearning_model.pkl", "wb") as f:
25         pickle.dump(Q_table, f)
26     print("[INFO] Q-table saved to qlearning_model.pkl")
27
28 def q_learning_move(game, player):
29     global Q_table, last_state, last_action, EPSILON, game_counter
30
31     current_state = state_str(game, player)
32     available_moves = game.available_moves()
33
34     if current_state not in Q_table:
35         Q_table[current_state] = {move: 0.0 for move in available_moves}
36

```



```

37     if random.random() < EPSILON:
38         center = 4
39         if center in available_moves and random.random() < 0.7:
40             action = center
41         else:
42             action = random.choice(available_moves)
43     else:
44         action = max(available_moves, key=lambda a: Q_table[current_state].get(a, 0.0))
45
46     if last_state is not None and last_action is not None:
47         future_q = max(Q_table[current_state].values()) if Q_table[current_state] else
0.0
48         old_q = Q_table[last_state].get(last_action, 0.0)
49         reward = 0
50         Q_table[last_state][last_action] = old_q + ALPHA * (reward + GAMMA * future_q -
old_q)
51
52     last_state = current_state
53     last_action = action
54     EPSILON = max(EPSILON * EPSILON_DECAY, EPSILON_MIN)
55
56     game_counter += 1
57     if game_counter % SAVE_FREQUENCY == 0:
58         save_Q_table_to_disk()
59
60     return action
61
62 def update_terminal(reward):
63     global Q_table, last_state, last_action
64     if last_state is not None and last_action is not None:
65         old_q = Q_table[last_state].get(last_action, 0.0)
66         Q_table[last_state][last_action] = old_q + ALPHA * (reward - old_q)
67     reset_episode()
68
69 def reset_episode():
70     global last_state, last_action
71     last_state = None
72     last_action = None
73
74 def save_model(filename="qlearning_model.pkl"):
75     with open(filename, "wb") as f:
76         pickle.dump(Q_table, f)
77     print(f"[INFO] Q-learning model saved to {filename}")
78
79 def load_model(filename="qlearning_model.pkl"):
80     global Q_table
81     try:
82         with open(filename, "rb") as f:
83             Q_table = pickle.load(f)
84             print(f"[INFO] Q-learning model loaded from {filename}")
85     except FileNotFoundError:
86         print("[WARN] No saved Q-table found. Starting fresh.")

```

## game.py

```

1 class TicTacToe:
2     def __init__(self):

```

```

3     self.board = [' '] * 9
4     self.current_winner = None
5     def print_board(self):
6         for row_idx in range(3):
7             row = self.board[row_idx * 3:(row_idx + 1) * 3]
8             indices = [str(i) for i in range(row_idx * 3, (row_idx + 1) * 3)]
9             print('| ' + ' | '.join(row) + ' |' + " <- " + ' '.join(indices))
10        print()
11
12    def available_moves(self):
13        return [i for i, spot in enumerate(self.board) if spot == ' ']
14
15    def empty_squares(self):
16        return ' ' in self.board
17
18    def make_move(self, square, letter):
19        if self.board[square] == ' ':
20            self.board[square] = letter
21            if self.check_winner(square, letter):
22                self.current_winner = letter
23            return True
24        return False
25
26    def check_winner(self, square, letter):
27        row_idx = square // 3
28        row = self.board[row_idx * 3:(row_idx + 1) * 3]
29        if all(spot == letter for spot in row):
30            return True
31
32        col_idx = square % 3
33        column = [self.board[col_idx + i * 3] for i in range(3)]
34        if all(spot == letter for spot in column):
35            return True
36
37        if square % 2 == 0:
38            diagonal1 = [self.board[i] for i in [0, 4, 8]]
39            if all(spot == letter for spot in diagonal1):
40                return True
41            diagonal2 = [self.board[i] for i in [2, 4, 6]]
42            if all(spot == letter for spot in diagonal2):
43                return True
44
45        return False

```

## main.py

```

1 import random
2 import os
3 import csv
4 import time
5 from datetime import datetime
6 import matplotlib.pyplot as plt
7
8 from game import TicTacToe
9 from algorithms import minimax, qlearning, baseline
10
11 def select_alpha_beta():

```

```

12     response = input("Use alpha-beta pruning for minimax? (y/n): ").strip().lower()
13     return response == 'y'
14
15 def get_move(game, player_letter, algorithm, use_alpha_beta, time_limit=30):
16     start_time = time.time()
17
18     if algorithm == "baseline":
19         move = baseline.baseline_move(game, player_letter)
20     elif algorithm == "minimax":
21         if use_alpha_beta:
22             move = minimax.minimax(game, player_letter, -float('inf'), float('inf'))["position"]
23         else:
24             move = minimax.minimax_no_ab(game, player_letter)["position"]
25     elif algorithm == "qlearning":
26         move = qlearning.q_learning_move(game, player_letter)
27     else:
28         move = random.choice(game.available_moves())
29
30     elapsed_time = time.time() - start_time
31     return move, elapsed_time
32
33 def play_game_matchup(matchup, use_alpha_beta):
34     game = TicTacToe()
35     if matchup == "1":
36         algo1, algo2 = "baseline", "minimax"
37     elif matchup == "2":
38         algo1, algo2 = "baseline", "qlearning"
39     elif matchup == "3":
40         algo1, algo2 = "minimax", "qlearning"
41     elif matchup == "4":
42         algo1, algo2 = "qlearning", "minimax"
43     else:
44         algo1, algo2 = "baseline", "minimax"
45
46     player1_letter, player2_letter = 'X', 'O'
47     turn = "player1" if random.random() < 0.5 else "player2"
48
49     moves_count = 0
50     algo1_total_time = 0
51     algo2_total_time = 0
52     algo1_moves = 0
53     algo2_moves = 0
54
55     while game.empty_squares():
56         moves_count += 1
57
58         if turn == "player1":
59             move, move_time = get_move(game, player1_letter, algo1, use_alpha_beta)
60             algo1_total_time += move_time
61             algo1_moves += 1
62
63         game.make_move(move, player1_letter)
64         if game.current_winner == player1_letter:
65             if algo1 == "qlearning": qlearning.update_terminal(10)
66             if algo2 == "qlearning": qlearning.update_terminal(-10)
67
68         algo1_avg_time = algo1_total_time / algo1_moves if algo1_moves > 0 else
0

```

```

69         algo2_avg_time = algo2_total_time / algo2_moves if algo2_moves > 0 else
0
70
71         return "player1", algo1, algo2, algo1_avg_time, algo2_avg_time,
moves_count
72         turn = "player2"
73     else:
74         move, move_time = get_move(game, player2_letter, algo2, use_alpha_beta)
75         algo2_total_time += move_time
76         algo2_moves += 1
77
78         game.make_move(move, player2_letter)
79         if game.current_winner == player2_letter:
80             if algo2 == "qlearning": qlearning.update_terminal(10)
81             if algo1 == "qlearning": qlearning.update_terminal(-10)
82
83             algo1_avg_time = algo1_total_time / algo1_moves if algo1_moves > 0 else
0
84             algo2_avg_time = algo2_total_time / algo2_moves if algo2_moves > 0 else
0
85
86         return "player2", algo1, algo2, algo1_avg_time, algo2_avg_time,
moves_count
87         turn = "player1"
88
89     if algo1 == "qlearning": qlearning.update_terminal(0)
90     if algo2 == "qlearning": qlearning.update_terminal(0)
91
92     algo1_avg_time = algo1_total_time / algo1_moves if algo1_moves > 0 else 0
93     algo2_avg_time = algo2_total_time / algo2_moves if algo2_moves > 0 else 0
94
95     return "tie", algo1, algo2, algo1_avg_time, algo2_avg_time, moves_count

```